

# Designing for Scalability on AWS

---



**Danny Jessee**

AWS CERTIFIED DEVOPS ENGINEER - PROFESSIONAL

@dannyjessee



# Overview



Architecting scalable solutions

Launch Configurations and  
Launch Templates

Create a new Auto Scaling Group

Auto Scaling policies

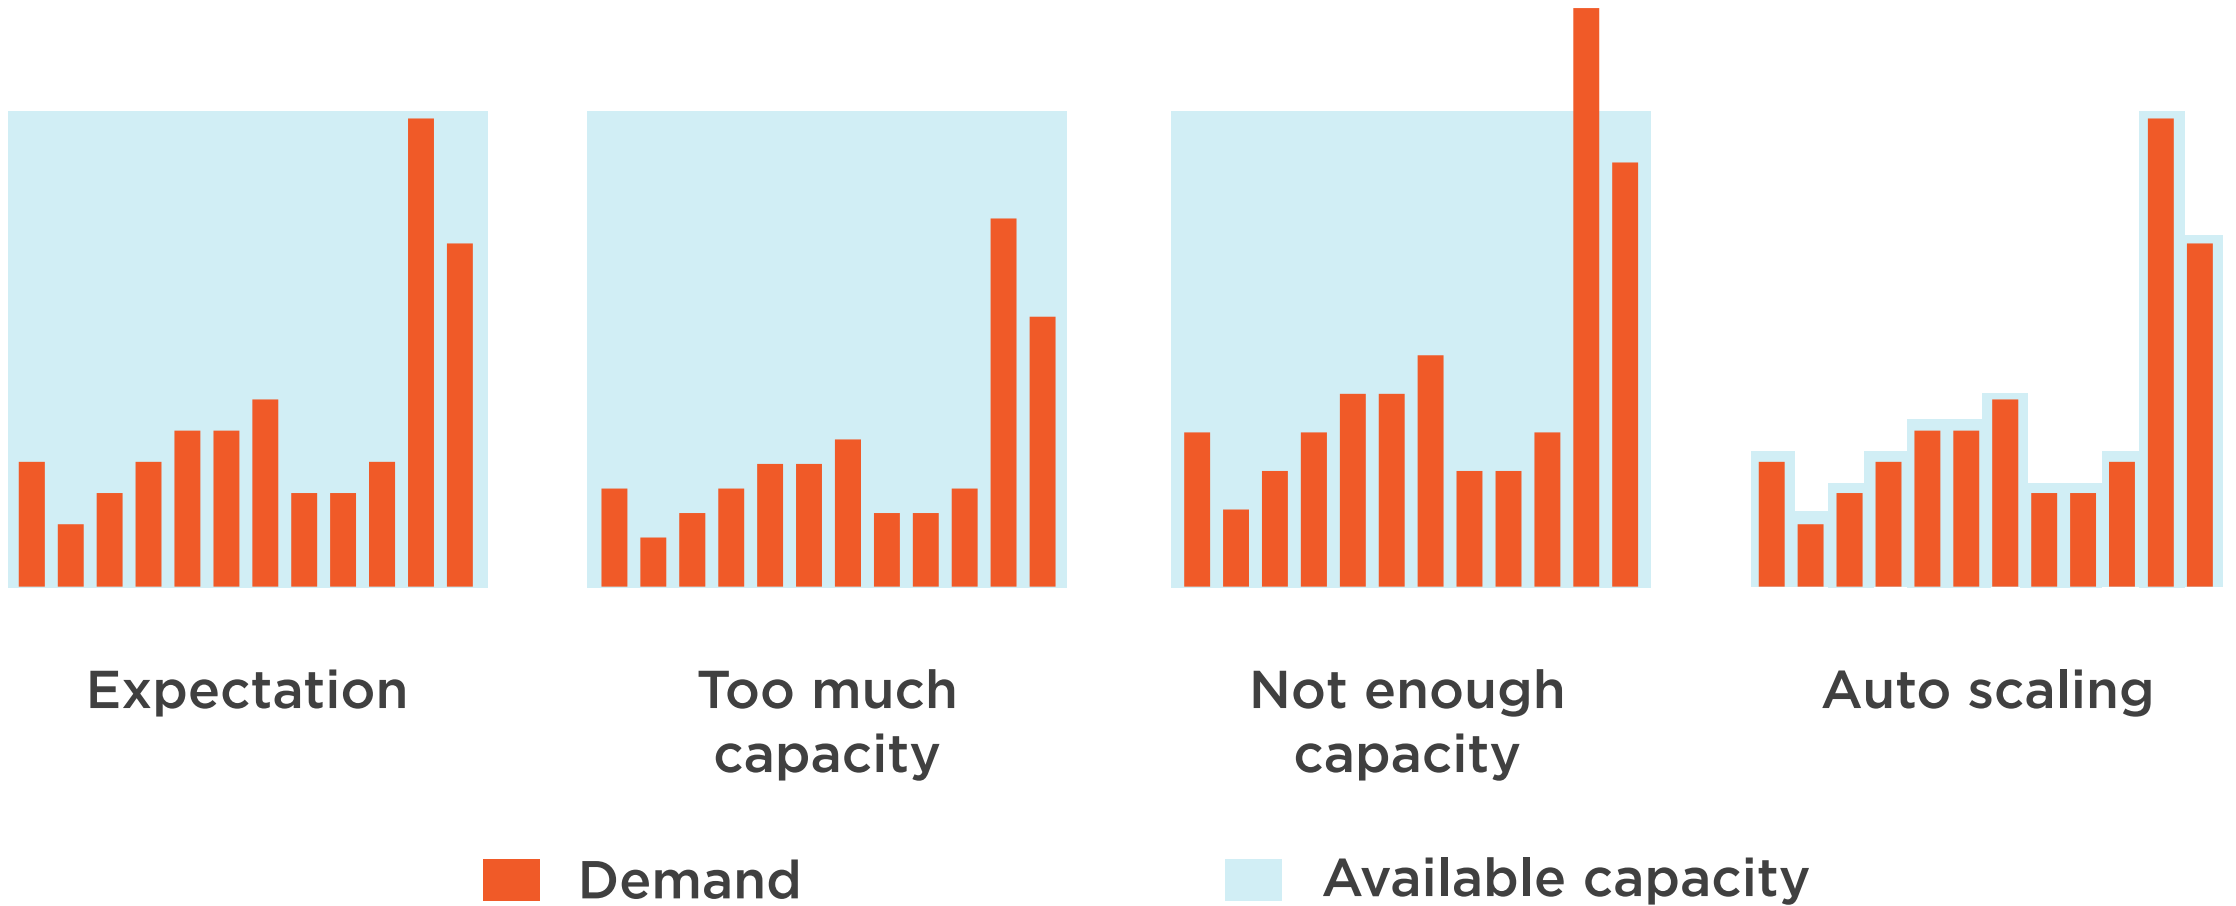
Auto Scaling lifecycle hooks

Scaling on a schedule


Self-healing nature of Auto  
Scaling Groups




# Capacity Planning



# Horizontal and Vertical Scaling



Horizontal scaling  
“Scaling out”




Vertical scaling  
“Scaling up”

# Launch Configurations



Can be defined from scratch



Can be created from a running EC2 instance



Can be reused across multiple Auto Scaling Groups



Are immutable



# Demo



Create a new Launch Configuration  
Spot instances and their benefits



# Launch Templates

**Recommended by AWS over  
Launch Configurations**

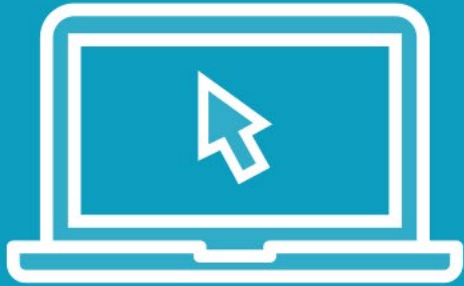
**Can be changed and have  
multiple versions**

**Allow for more advanced  
configuration options**

**Allow for mixing instance  
types and classes**



# Demo



Create a new Launch Template


Benefits of using Launch Templates

Create a Launch Template from a  
Launch Configuration







# Auto Scaling



Right-sizing  
capacity needs



Logical  
collection of  
instances

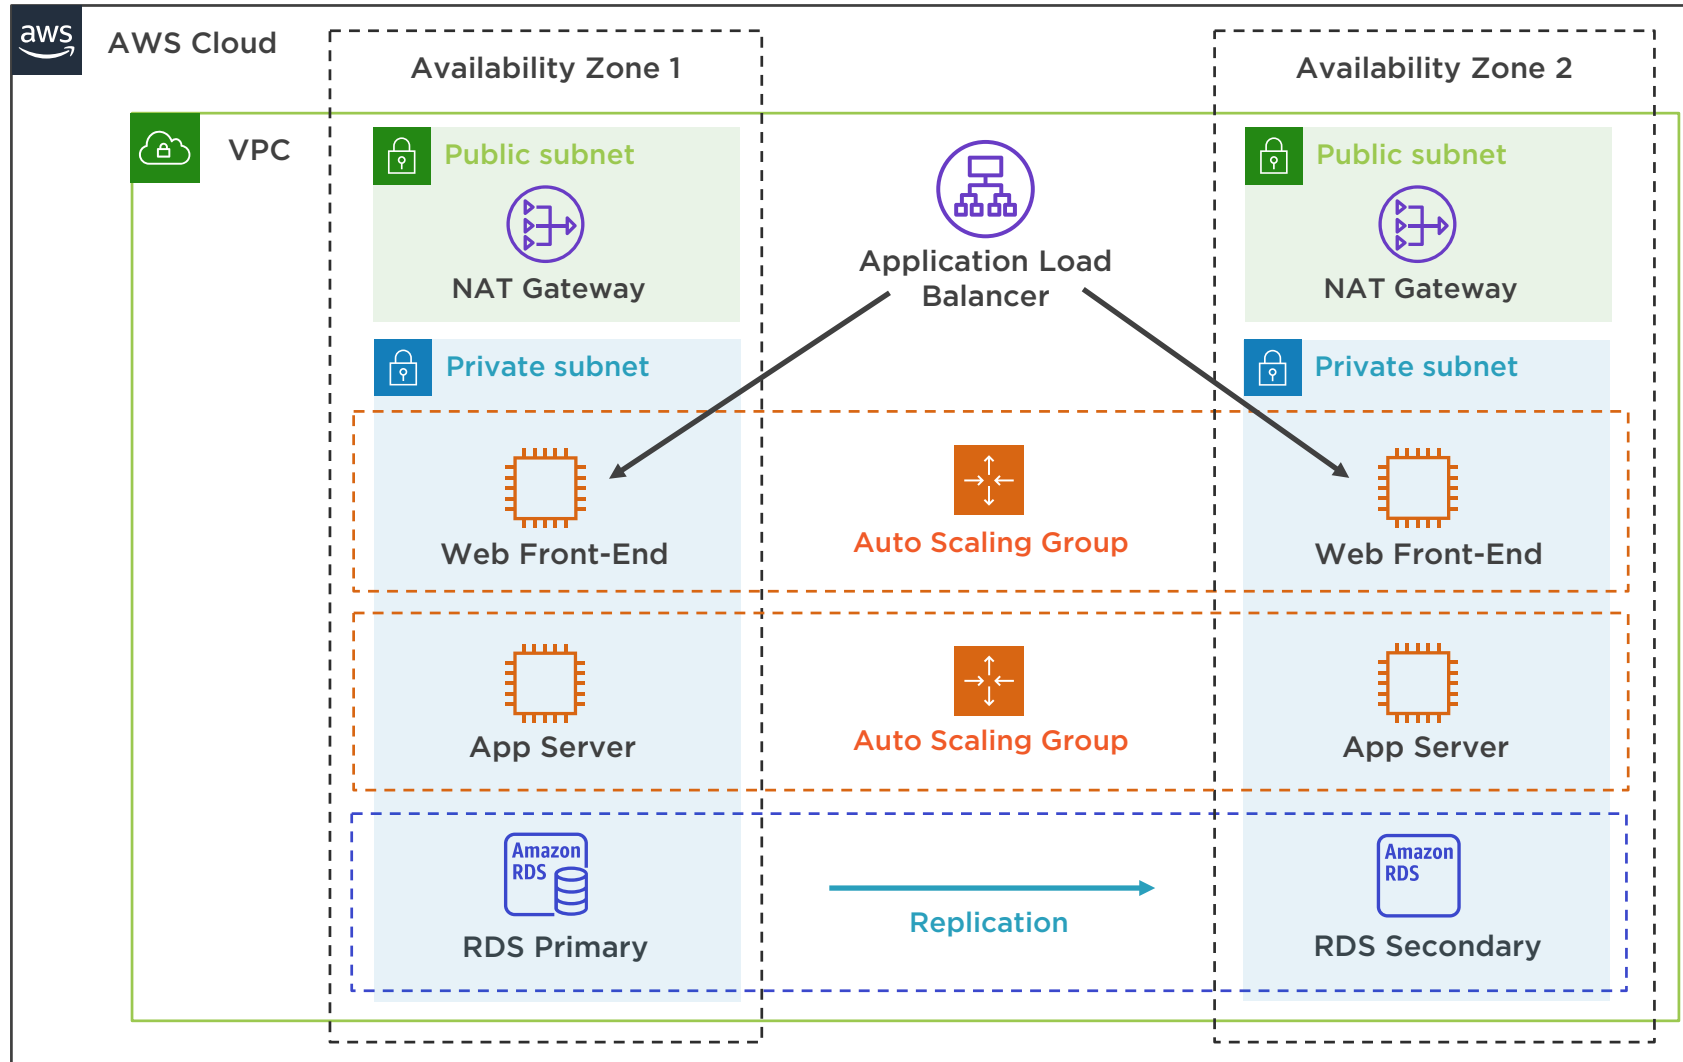


Spread across  
Availability  
Zones

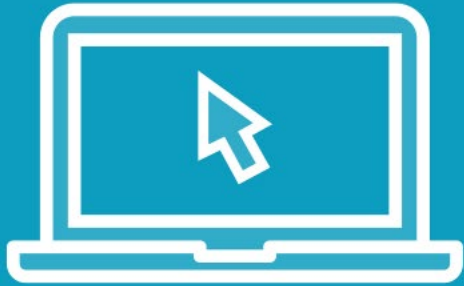


Performs health  
checks on  
instances

# Elastic Load Balancing and Auto Scaling



# Demo



Create a new Auto Scaling Group

Associate an Auto Scaling Group with an Application Load Balancer



# Auto Scaling Policies

Manually adjust minimum, maximum, or desired capacity values

Automatically adjust based on metric values




## Target Tracking

Auto Scaling will maintain target value



## Simple

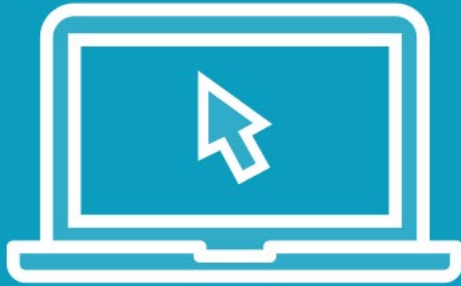
Auto Scaling acts based on CloudWatch Alarms



## Step Scaling

Auto Scaling acts based on banded value ranges

# Demo

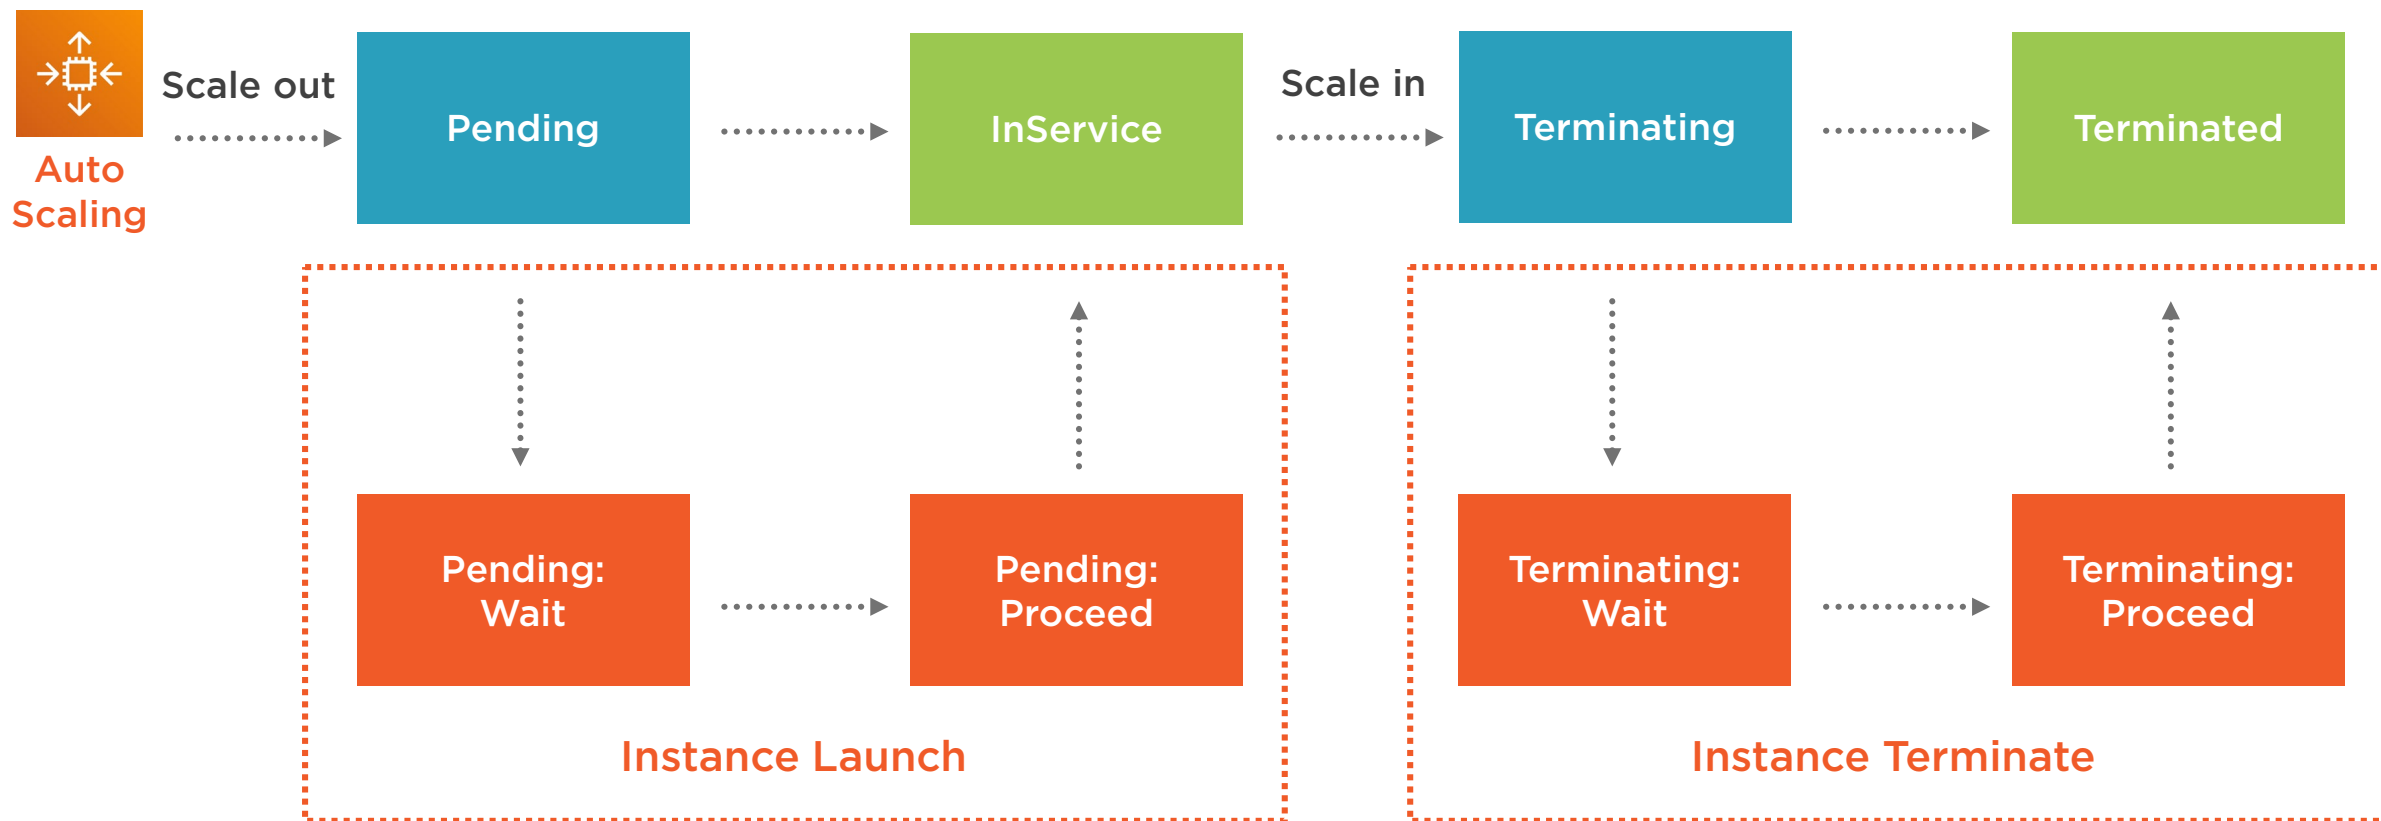


Configure scaling policies for an Auto Scaling Group

Target Tracking, Simple, and Step Scaling policies



# Lifecycle Hooks



# Demo




Create a lifecycle hook

CloudWatch Events rule and targets



# Dynamically Scaling on a Schedule



Set a fixed value for minimum, maximum, or desired capacity

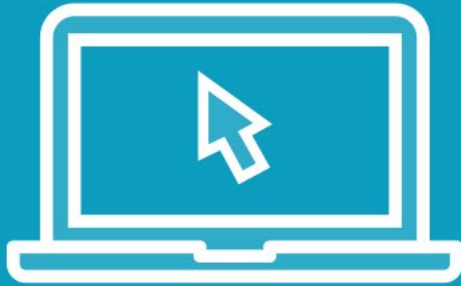
Can be one-time or recurring

Cron expressions for recurrence pattern

Start and end date



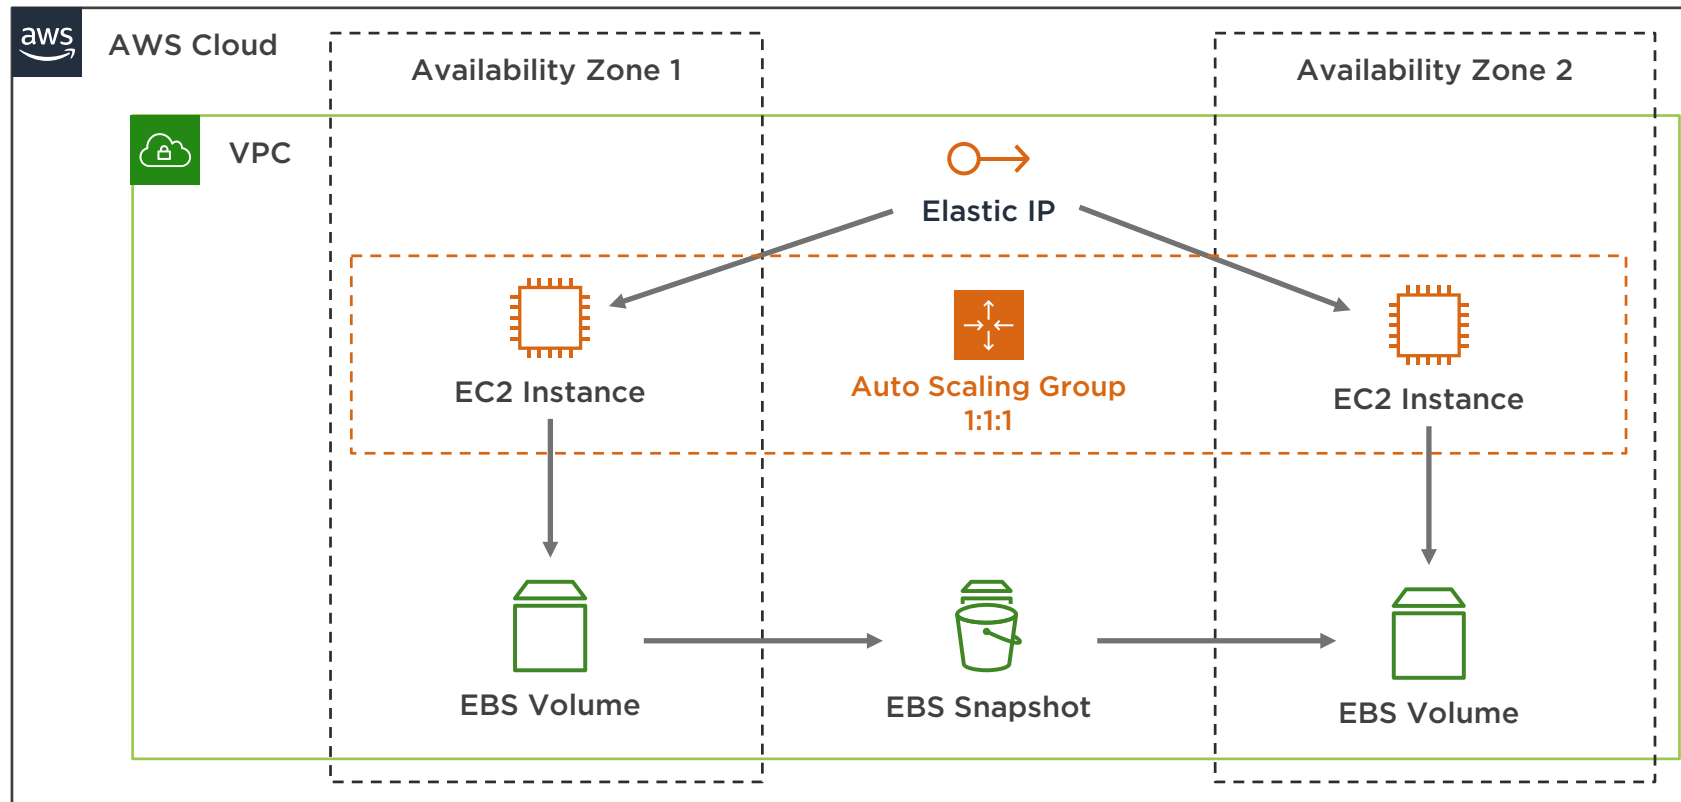
# Demo



## Create a Scheduled Action



# Self-Healing Instance with Auto Scaling



# Review



**AWS Auto Scaling**

**Launch Configurations and  
Launch Templates**

**Auto Scaling policies**

**Auto Scaling lifecycle hooks**

**Scheduled actions**

**Self-healing with Auto Scaling Groups**



Up Next:

# Designing Highly Available and Scalable Databases

---

