

Designing Highly Available and Scalable Databases on AWS



Danny Jessee

AWS CERTIFIED DEVOPS ENGINEER - PROFESSIONAL

@dannyjessee



Overview



Amazon Relational Database Service (RDS)

Supported database engines

RDS Multi-AZ deployments

Vertical and horizontal scaling

Amazon DynamoDB



3-tier Architecture




Front-end
web server

Application
server

Database
server

Requires ongoing patching and maintenance

Requires DBA labor to configure replication and failover




EC2 instance




Amazon Relational Database Service (RDS)



Fully-managed



Leverages EC2 instances
and EBS volumes




AWS handles patching
and maintenance



Automated backups for
point-in-time recovery



User-initiated snapshots



Multi-AZ deployments



Running Your Database on EC2



Shell access to server file system or other server-side dependencies



Willing to expend DBA labor in exchange for lower usage costs



Additional work required to support maintenance and failover scenarios

Hands-on lab: Build your own highly available database servers on EC2



Supported RDS Database Engines



MySQL, MariaDB, PostgreSQL

Oracle and Microsoft SQL Server

“Easy create” options

Oracle supports “bring your own license”
and “license included” pricing

Microsoft SQL Server “license included”



Amazon Aurora

MySQL and
PostgreSQL-
compatible

5x faster than
MySQL for
1/10 the cost

Up to 64 TB
per database
instance

Data written
6x across 3
availability zones

Up to 15 read
replicas across 3
availability zones

Global databases
can span multiple
AWS regions



RDS Multi-AZ Deployments



Highly available by spanning two availability zones



Can convert a non-Multi-AZ deployment



RDS maintains a synchronous standby replica



Automatic failover in the event of an outage or maintenance



Expose a single DNS endpoint




Demo



Create a new RDS Multi-AZ deployment
Failover scenarios



Scaling with RDS



Scale up
by changing
instance class



Add storage
or provisioned
IOPS



Scale out
by adding
read replicas



Handle increased
demand for
read operations



Demo



Change instance type and storage

Add an asynchronous read replica



Introduction to DynamoDB



DynamoDB Overview



Serverless

K	V

Schemaless
(NoSQL)



Not easy to
convert to/from
relational




Continuous or
on-demand
backups

DynamoDB Tables



No database, table is highest-level entity



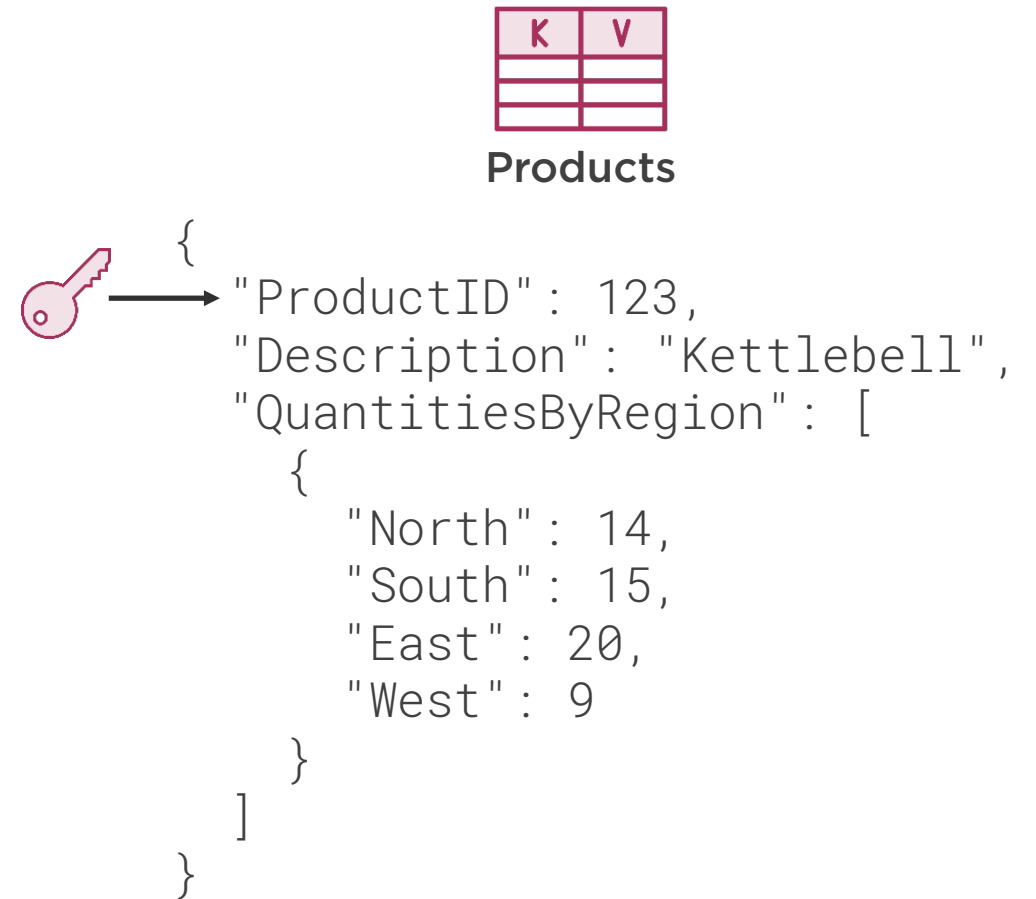

Essentially unlimited storage and scalability



Data automatically replicated across multiple availability zones



Relational Database vs. DynamoDB



DynamoDB Secondary Indexes

Local Secondary Indexes (LSI)

Only created when table is created

Same partition key, different sort key

Local to a partition

Can return non-projected attributes

Consume provisioned RCU/WCU
from parent table

Global Secondary Indexes (GSI)

Can be created after table is created

Different partition and sort keys

Global to a table

Can not return non-projected attributes

Separately provisioned RCU/WCU



DynamoDB Provisioned Capacity

Read/Write
capacity units
(RCU/WCU)

Throttling when you
exceed provisioned
capacity

Can use
auto scaling

Estimate average
size of data read
and write requests

Estimate number
of reads and writes
per second

Cost is based on
provisioned
RCU and WCU



DynamoDB Provisioned Capacity

Read
Capacity Units
(RCU)

For items up to 4 KB in size:

1 RCU = 2 eventually consistent read requests per second

1 RCU = 1 strongly consistent read request per second

2 RCU = 1 transactional read request per second



DynamoDB Provisioned Capacity

Write
Capacity Units
(WCU)

For items up to 1 KB in size:

**1 WCU = 1 standard write request
per second**

**2 WCU = 1 transactional write request
per second**



Amazon DynamoDB Accelerator (DAX)



Caching layer to
prevent duplicate
read requests



Fully managed,
in-memory cache



Provision a DAX
cluster in front of your
DynamoDB table




DynamoDB Streams and Global Tables




DynamoDB writes data across AZs by default




Streams can allow data to be written across regions




Streams capture all activity within a table



Can include before and after images of items

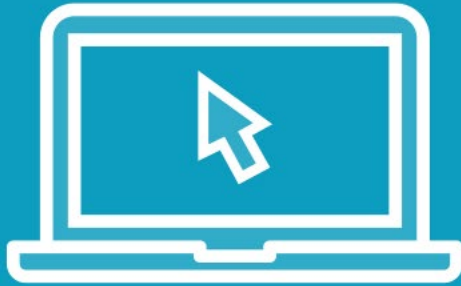


Global Tables automate replication across regions



Global, fault-tolerant applications

Demo



Create a new DynamoDB table

Partition and sort keys

Secondary indexes

Provision read and write capacity

DynamoDB Global Tables



Review



Amazon RDS

Amazon Aurora

RDS Multi-AZ deployments

Read replica instances

Amazon DynamoDB

DynamoDB Streams and Global Tables



Up Next:

Disaster Recovery Strategies on AWS

